

Oliver Grabner

oliver.charles.rene.grabner@gmail.com | olivergrabner.com | linkedin.com/in/olivergrabner | github.com/OliverGrabner

EDUCATION

TEXAS A&M UNIVERSITY

Bachelor of Science in Computer Science

- Goldman Sachs Excellence Scholarship

Expected May 2027

GPA - 4.00

EXPERIENCE

Google

Software Engineer Intern

- Incoming Software Engineer Intern

May 2026 – August 2026

Seattle, WA

Samsung Semiconductor

Software Engineer Intern

- Reduced API latency by 63% by designing 5 asynchronous FastAPI endpoints for concurrent processing of large-scale FAB data from secure AWS S3 data lake, enabling real-time analytics
- Developed a full-stack Streamlit dashboard to automate system access audits, reducing weekly manual work by 97% (2 hours to 3 minutes) and ensuring 100% compliance by eliminating manual errors
- Engineered an AI Agent (Llama 3 70B) with a LangChain router that translates user queries into Python code executed in a secure sandbox for data analysis on 9 data sources or into function calls for workflow automation
- Architected a reusable RAG platform to deploy specialized workflow chatbots, integrating query rewriting and hybrid search (BM25 and FAISS) to deliver relevant, domain-specific answers

May 2025 – August 2025

Austin, TX

Undergraduate Researcher

Texas A&M University - InfoLab

- Implemented a chunk segmentation pipeline for feeding long-form transcripts into LLMs; boosted F1 performance by 18% and halved variance on 1-7B parameter models in disfluency removal tasks
- Built an LLM-as-judge evaluation framework using Llama and Qwen to audit disfluency removal outputs, comparing LLM judgments against human annotations to quantify agreement in speech evaluation

September 2024 – Present

College Station, TX

AncientScript

Data Science Intern

- Built a torchvision augmentation pipeline expanding training data 3x and a polyline-to-tensor preprocessing function for a U-Net model achieving 94% realism on historical handwritten records

June 2024 – August 2024

Austin, TX

PROJECTS

Composer | [Link](#) | *Python, FastAPI, Gemini, MCP, SSE, WebSocket, PostgreSQL*

- Built an 8-tool autonomous agent loop that generates typed architecture diagrams with API contracts and entity schemas; streams decisions via a thread-to-async bridge over SSE
- Exposed an 11-tool MCP server enabling external AI tools to read/write diagrams with real-time WebSocket sync

Stock Wordle | [Link](#) | *Kubernetes, Docker, Java Spring Boot, Python, PostgreSQL*

- Deployed a Kubernetes cluster hosting a 2-replica Spring Boot API, PostgreSQL with PersistentVolumes, and Python CronJobs for daily batch processing; frontend hosted on Vercel
- Automated Python data pipeline that fetches from external APIs, performing UPSERTs with 30-day lookback

Pokemon Card Generator | [Link](#) | *Python, FastAPI, PyTorch, PostgreSQL, Docker*

- Trained a Deep Convolutional GAN (DCGAN) in PyTorch on 11,000+ images to generate novel 96x64 card art
- Developed an async FastAPI endpoint for real-time GAN inference, serving in-memory PyTorch tensors as base64

PUBLICATIONS

M. Teleki, S. Janjur, H. Liu, **O. Grabner**, et al. "Z-Scores: A Metric for Linguistically Assessing Disfluency Removal." *ICASSP 2026*

M. Teleki, S. Janjur, H. Liu, **O. Grabner**, et al. "DRES: Benchmarking LLMs for Disfluency Removal." *arXiv 2025*

TECHNICAL SKILLS

Languages: Python, C++, Java, Go, SQL, JavaScript, Assembly

Infrastructure & Tools: Kubernetes, Docker, Linux, Git, AWS S3, PostgreSQL, Vercel, asyncio, Maven

Frameworks: Spring Boot, Next.js, React, FastAPI, Django, PyTorch, LangChain, Hugging Face, OpenAI API