

Oliver Grabner

📞 (669) 271-9070 | 📩 olivergrabner.cs@gmail.com | 💼 linkedin.com/in/olivergrabner | 🐾 github.com/OliverGrabner

EDUCATION

TEXAS A&M UNIVERSITY

Bachelor of Science in Computer Science

- Goldman Sachs Excellence Scholarship

Expected May 2027

GPA - 4.00

EXPERIENCE

Samsung Semiconductor

May 2025 – August 2025

Austin, TX

Software Engineer Intern

- Reduced API latency by 63% by designing 5 asynchronous FastAPI endpoints for concurrent processing of large-scale FAB data from secure AWS S3 data lake, enabling real-time analytics
- Developed a full-stack Streamlit dashboard to automate system access audits, reducing weekly manual work by 97% (2 hours to 3 minutes) and ensuring 100% compliance by eliminating manual errors
- Engineered an AI Agent (Llama 3 70B) with a LangChain router that translates user queries into Python code executed in a secure sandbox for data analysis on 9 data sources or into function calls for workflow automation
- Architected a reusable RAG platform to deploy specialized workflow chatbots, integrating advanced query rewriting and hybrid search (BM25 and FAISS) to deliver relevant, domain-specific answers

Undergraduate Researcher

September 2024 – Present

College Station, TX

Texas A&M University - InfoLab

- Researching and developing a multi-modal recommender system fusing speech, LLM, and collaborative signals
- Implemented a chunk segmentation pipeline for feeding long-form transcripts into LLMs; boosted F1 performance by 18% and halved variance on 1-7B parameter models in disfluency removal tasks
- Automated reproducible benchmarking for 11 open-source LLMs (1B - 70B parameters) using HuggingFace + Python pipelines to log ROUGE and BERTScore for across k-shot evaluations

AncientScript

June 2024 – August 2024

Cedar Park, TX

Data Science Intern

- Developed a comprehensive torchvision augmentation pipeline for historical handwritten records, expanding the effective training set 3x and improving test accuracy 4 pp by discouraging spurious correlations
- Built a data processing function to transform annotated polyline underlines into normalized tensors, feeding a U-Net model the training data required to reach a 94% realism score

PUBLICATIONS

M. Teleki, S. Janjur, H. Liu, **O. Grabner**, et al. "Z-Scores: A Metric for Linguistically Assessing Disfluency Removal." *IEEE ICASSP 2026* (Under Review)

M. Teleki, S. Janjur, H. Liu, **O. Grabner**, et al. "DRES: Benchmarking LLMs for Disfluency Removal." *IEEE ICASSP 2026* (Under Review)

PROJECTS

Boba Shop POS System | Java, Swing, JDBC, PostgreSQL, AWS

- Developed a full-stack Java POS system with a Swing GUI client and a cloud-hosted AWS PostgreSQL backend
- Scripted the generation of \$5M+ in realistic sales data to test and populate a relational PostgreSQL schema
- Implemented a manager dashboard, authoring 15+ SQL queries to generate reports on sales trends and inventory

Pokemon Card Generator | Link | Python, FastAPI, PyTorch, PostgreSQL, Docker, Vercel, Render

- Trained a Deep Convolutional GAN (DCGAN) in PyTorch on 11,000+ images to generate novel 96x64 card art
- Developed an async FastAPI endpoint for real-time GAN inference, serving in-memory PyTorch tensors as base64
- Built a PostgreSQL-backed community gallery with user sharing, voting, and paginated sorting (popular/recent)
- Designed schema with custom SQL indexes to optimize sorting; deployed via split Vercel/Render (Docker) pipeline

TECHNICAL SKILLS

Languages: Python, C++, Java, SQL, JavaScript, HTML, CSS, Assembly

Frameworks & Libraries: FastAPI, Django, Flask, React, Vite, Tailwind CSS, Pandas, Matplotlib, Streamlit, PyQt5

Tools & Technologies: Docker, AWS S3, boto3, Git, Kubernetes, asyncio, Netlify, GA4, Confluence

Machine Learning: PyTorch, LangChain, HuggingFace, OpenCV, NumPy, torchvision, scikit-learn, OpenAI API