

Oliver Grabner

oliver.charles.rene.grabner@gmail.com | olivergrabner.com | linkedin.com/in/olivergrabner | github.com/OliverGrabner

EDUCATION

Texas A&M University

August 2023 – May 2027

Bachelor of Science in Computer Science

GPA: 4.00

- Goldman Sachs Excellence Scholarship

EXPERIENCE

Google

May 2026 – August 2026

Software Engineer Intern

Kirkland, WA

- Built an async Kotlin RPC client for **Take Notes For Me**, Google Meet's AI note-taking feature; reduced mean end-to-end latency **34% (751 to 499 ms)** versus the legacy client in benchmarks
- Diagnosed and fixed a protobuf **oneof** collision in the document delivery integration, replacing a workaround that relied on deprecated fields with a supported code path
- Built a fluent Python library for PM, research, and UX prototyping, eliminating manual character-index tracking in protobuf documents; migrated **12 active prototypes**
- Authored two **LLM agent skills** with eval suites, cutting design iteration per task **9.5x (262 to 27.5 s)** and token usage **79% (590k to 121k)**; task success rate doubled (24% to 48%)
- Migrated Take Notes For Me headers and weights from Roboto to **Google Sans Flex**; **launched to 100% GA**

Samsung Semiconductor

May 2025 – August 2025

Software Engineer Intern

Austin, TX

- **Reduced API latency 63%** across 5 FastAPI endpoints, overlapping concurrent S3 fetches with computation and offloading blocking calls to thread pools to keep the event loop unblocked
- Built a Llama 3 70B **AI agent** whose LangChain router sends each query down one of two paths: sandboxed Python for analysis across 9 data sources, or function calls for workflow automation

Texas A&M University

September 2024 – Present

Undergraduate Researcher - InfoLab

College Station, TX

- Implemented a chunk segmentation pipeline for feeding long-form transcripts into LLMs; **raised F1 by 18 points** and halved run-to-run variance on 1-7B parameter models in disfluency removal tasks
- Extended **mini-SWE-agent** with structured planning and tool-observation pruning; evaluated inference-cost and task-success tradeoffs on **SWE-bench Verified** using Qwen3.5-27B/vLLM

PROJECTS

Composer | [Link](#) | Python, FastAPI, Gemini, MCP, React, WebSocket, PostgreSQL

- Built a FastAPI application that converts natural language into editable architecture diagrams; exposed MCP tools for external agents, **serving 4,000+ tool calls** since launch
- Implemented a thread-to-async bridge for SSE streaming and agent suspension/resumption across HTTP requests, with WebSocket synchronization for live diagram edits

Stock Wordle | [Link](#) | Java, Spring Boot, PostgreSQL, Kubernetes, Python

- Built a **Java/Spring Boot REST API** for a daily stock-guessing game, implementing guess evaluation, hints, and price-history responses with Spring Data JPA and PostgreSQL
- Deployed a two-replica API and persistent PostgreSQL storage on Kubernetes; scheduled Python CronJobs to ingest market data and update daily puzzles through SQL UPSERTs

Pokemon Card Generator | [Link](#) | Python, FastAPI, PyTorch, PostgreSQL, Docker

- Trained a PyTorch DCGAN on **11,000 card images** and integrated it into a FastAPI service for on-demand generation, with PostgreSQL-backed sharing, voting, and paginated galleries

SELECTED PUBLICATIONS

M. Teleki, S. Janjur, H. Liu, **O. Grabner**, et al. “Z-Scores: A Metric for Linguistically Assessing Disfluency Removal.” *ICASSP 2026*

M. Teleki, S. Janjur, H. Liu, **O. Grabner**, et al. “Conversational Speech Reveals Structural Robustness Failures in SpeechLLM Backbones.” *Under review, NeurIPS 2026 (Evaluations & Datasets)*

TECHNICAL SKILLS

Languages: Python, Kotlin, Java, C++, SQL, JavaScript, TypeScript, Protocol Buffers

Frameworks: Spring Boot, FastAPI, React, Next.js, PyTorch, LangChain, asyncio, JUnit/Mockito

Infrastructure: Kubernetes, Docker, PostgreSQL, gRPC, AWS S3, Vercel, Git, Linux